

# Transaction Mempool

## Objectives

1. To provide necessary methods to **create and close mempool**
2. To provide necessary methods to **add or drop transactions** to/from the mempool
3. To provide necessary methods to **validate transactions** and decide whether to kick them or not.
4. To provide necessary methods to help block proposer (will be searchers in MEV-Boost case) to **compose a transaction batch from the mempool**


## High Level Overview

- Mempool is a place where pending transactions stay before they are executed.
- Creating a mempool means **creating a pool structure, maintaining a list of transactions**, and **running a goroutine** that subscribes the addition of new block and updates the state of mempool
- After EIP-4844, there comes "blob transaction", whose syntax and validation logic is quite different from legacy transactions.
- So the mempool should also be separated from the legacy mempool.
- For this purpose, geth enables the client to create various transaction-specific pools, called `subpool`.
- Subpools share the same interface and managed by primary transaction pool (`txpool.go`), since they have to be updated and assembled into one coherent view for block production.
- Every pool **MUST** implement all the methods in `subpool.go` ⇒ Every pool is an instance of subpool interface.

Currently two pools are implemented like this way: `blobpool` and `legacypool`. Any additional tx type can generate extra txpool.

For instance, there will be `aapool` if native account abstraction is implemented on Ethereum.

(Look at [core/txpool](#) )



## Code Review

- Let's take a look at subpool interface first.

[core/txpool/subpool.go](#)

```
type SubPool interface {
    // Filter is a selector used to decide whether a transaction should be added
    // to this particular subpool.
    Filter(tx *types.Transaction) bool

    // Init sets the base parameters of the subpool, allowing it to load any saved
    // transactions from disk and also permitting internal maintenance routines to
    // start up.
    //
```

```
// These should not be passed as a constructor argument - nor should the pool
// start by themselves - in order to keep multiple subpools in lockstep with
// one another.
```

```
Init(gasTip *big.Int, head *types.Header, reserve AddressReserver) error
```

```
// Close terminates any background processing threads and releases any held
// resources.
```

```
Close() error
```

```
// Reset retrieves the current state of the blockchain and ensures the content
// of the transaction pool is valid with regard to the chain state.
```

```
Reset(oldHead, newHead *types.Header)
```

```
// SetGasTip updates the minimum price required by the subpool for a new
// transaction, and drops all transactions below this threshold.
```

```
SetGasTip(tip *big.Int)
```

```
// Has returns an indicator whether subpool has a transaction cached with the
// given hash.
```

```
Has(hash common.Hash) bool
```

```
// Get returns a transaction if it is contained in the pool, or nil otherwise.
```

```
Get(hash common.Hash) *types.Transaction
```

```
// Add enqueues a batch of transactions into the pool if they are valid. Due
// to the large transaction churn, add may postpone fully integrating the tx
// to a later point to batch multiple ones together.
```

```
Add(txs []*types.Transaction, local bool, sync bool) []error
```

```
// Pending retrieves all currently processable transactions, grouped by origin
// account and sorted by nonce.
```

```
Pending(enforceTips bool) map[common.Address][]*LazyTransaction
```

```
// SubscribeTransactions subscribes to new transaction events. The subscribe
// can decide whether to receive notifications only for newly seen transactions
// or also for reorged out ones.
```

```
SubscribeTransactions(ch chan← core.NewTxsEvent, reorgs bool) event.Subs
```

```
// Nonce returns the next nonce of an account, with all transactions executable  
// by the pool already applied on top.
```

```
Nonce(addr common.Address) uint64
```

```
// Stats retrieves the current pool stats, namely the number of pending and the  
// number of queued (non-executable) transactions.
```

```
Stats() (int, int)
```

```
// Content retrieves the data content of the transaction pool, returning all the  
// pending as well as queued transactions, grouped by account and sorted by  
Content() (map[common.Address][]*types.Transaction, map[common.Address]
```

```
// ContentFrom retrieves the data content of the transaction pool, returning the  
// pending as well as queued transactions of this address, grouped by nonce.
```

```
ContentFrom(addr common.Address) ([]*types.Transaction, []*types.Transacti
```

```
// Locals retrieves the accounts currently considered local by the pool.
```

```
Locals() []common.Address
```

```
// Status returns the known status (unknown/pending/queued) of a transaction  
// identified by their hashes.
```

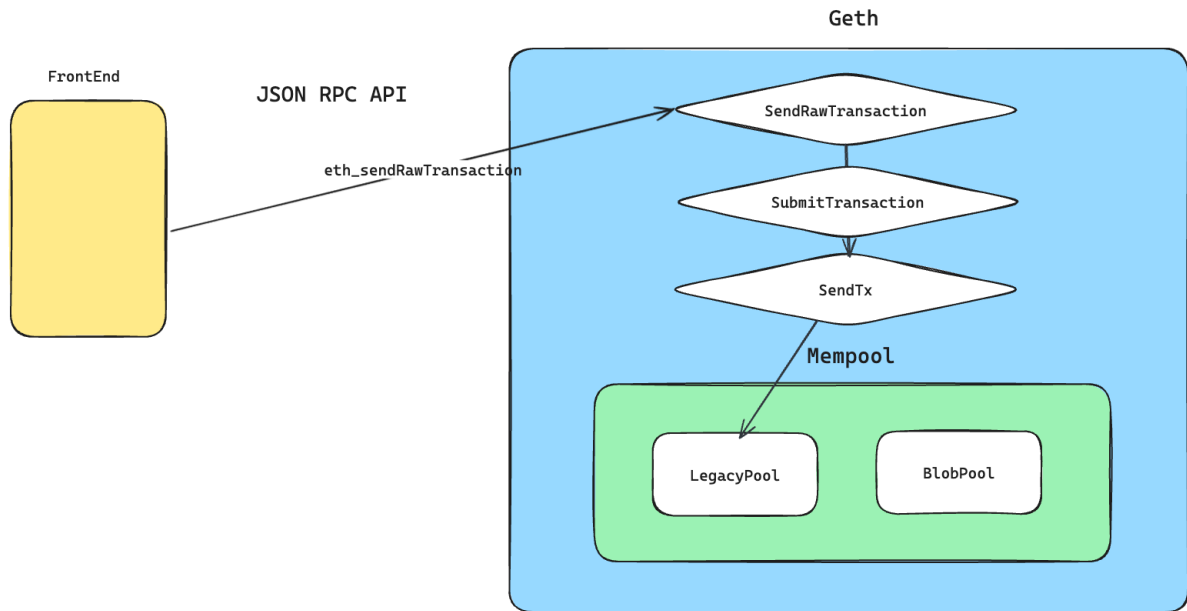
```
Status(hash common.Hash) TxStatus
```

```
}
```

- There are methods that creates / closes subpool, add / filter transactions inside the subpool, and retrieve status of transactions or the subpool.
- I'll introduce two of widely-used functions, `Add` & `Reset`.

#### `Add` & `Reset`

1. `Add(txs []*types.Transaction, local bool, sync bool) []error`



- This is a function that adds a set of transactions to the subpool.

▼ code

- When you send tx through JSON RPC API of Infura, Alchemy, or whatever, you are calling `eth_sendTransaction` or `eth_sendRawTransaction` method.
- Both methods are at [internal/ethapi/api.go](https://github.com/ethereum/go-ethereum/blob/master/internal/ethapi/api.go), and they are implemented like the following code.

```

func (s *TransactionAPI) SendRawTransaction(ctx context.Context, in
    tx := new(types.Transaction)
    if err := tx.UnmarshalBinary(input); err != nil {
        return common.Hash{}, err
    }
    return SubmitTransaction(ctx, s.b, tx)
}

```

and they calls `SubmitTransaction` method, which triggers `SendTx` method at [eth/api\\_backend.go](https://github.com/ethereum/go-ethereum/blob/master/internal/ethapi/api.go).

```
func (b *EthAPIBackend) SendTx(ctx context.Context, signedTx *types.Transaction) error {
    return b.eth.txPool.Add([]*types.Transaction{signedTx}, true, false)
}
```

- Add function is implemented like the following code in legacy pool.

```
func (pool *LegacyPool) Add(txs []*types.Transaction, local, sync bool) error {
    // Filter out known ones without obtaining the pool lock or recovering
    var (
        errs = make([]error, len(txs))
        news = make([]*types.Transaction, 0, len(txs))
    )
    for i, tx := range txs {
        // If the transaction is known, pre-set the error slot
        if pool.all.Get(tx.Hash()) != nil {
            errs[i] = ErrAlreadyKnown
            knownTxMeter.Mark(1)
            continue
        }
        // Exclude transactions with basic errors, e.g invalid signatures and
        // insufficient intrinsic gas as soon as possible and cache sender
        // in transactions before obtaining lock
        if err := pool.validateTxBasics(tx, local); err != nil {
            errs[i] = err
            invalidTxMeter.Mark(1)
            continue
        }
        // Accumulate all unknown transactions for deeper processing
        news = append(news, tx)
    }
    if len(news) == 0 {
        return errs
    }
}
```

```

// Process all the new transaction and merge any errors into the ori
pool.mu.Lock()
newErrs, dirtyAddrs := pool.addTxsWithLocked(news, local)
pool.mu.Unlock()

var nilSlot = 0
for _, err := range newErrs {
    for errs[nilSlot] != nil {
        nilSlot++
    }
    errs[nilSlot] = err
    nilSlot++
}
// Reorg the pool internals if needed and return
done := pool.requestPromoteExecutables(dirtyAddrs)
if sync {
    <-done
}
return errs
}

```

- It executes a for loop for every transactions.
- First, it validates if the transaction meets all requirements (valid signature, sufficient gas, ...)
- And it uses mutex to avoid errors from concurrency:
- Inside the mutex lock and unlock, `addTxsWithLocked` function adds the transaction into the pool. Since using mutex means that there is somewhere using goroutine inside the `addTxsWithLocked` function.

```

func (pool *LegacyPool) addTxsWithLocked(txs []*types.Transaction,
    dirty := newAccountSet(pool.signer)
    errs := make([]error, len(txs))
    for i, tx := range txs {

```

```

    replaced, err := pool.add(tx, local)
    errs[i] = err
    if err == nil && !replaced {
        dirty.addTx(tx)
    }
}
validTxMeter.Mark(int64(len(dirty.accounts)))
return errs, dirty
}

```

Unfortunately, I failed to find the exact place where the concurrency matters :(

IMO, maybe this line matters:


```
if uint64(pool.all.Slots()+numSlots(tx))...
```

This line checks the current slots of the subpool. If the pool is full, it drops the transaction. And, the slot data is a global state of each subpool. If each transaction is executed in a concurrent manner, checking in this line may produce unexpected error.

## 2. Reset(oldHead, newHead \*types.Header)

- Reset retrieves the current state of the blockchain and ensures the content of the transaction pool is valid with regard to the chain state.





- Think of a situation of reorg: Transactions in the 'orphan blocks' need to be back to the mempool.
- **Reset** updates the content of the pool with the given **newHead** which means a header of the new block, bringing all txs in the re-orged blocks.

▼ code

```

func (pool *LegacyPool) reset(oldHead, newHead *types.Header) {
    // If we're reorging an old state, reinject all dropped transactions
    var reinject types.Transactions

    if oldHead != nil && oldHead.Hash() != newHead.ParentHash {
        // If the reorg is too deep, avoid doing it (will happen during fast sync)
        oldNum := oldHead.Number.Uint64()
        newNum := newHead.Number.Uint64()

        if depth := uint64(math.Abs(float64(oldNum) - float64(newNum))); c
            log.Debug("Skipping deep transaction reorg", "depth", depth)
        } else {
            // Reorg seems shallow enough to pull in all transactions into mem
            var (
                rem = pool.chain.GetBlock(oldHead.Hash(), oldHead.Number.Ui

```

```

    add = pool.chain.GetBlock(newHead.Hash(), newHead.Number.
)
if rem == nil {
    // This can happen if a setHead is performed, where we simply
    // head from the chain.
    // If that is the case, we don't have the lost transactions anymore
    // there's nothing to add
    if newNum >= oldNum {
        // If we reorged to a same or higher number, then it's not a ca
        log.Warn("Transaction pool reset with missing old head",
            "old", oldHead.Hash(), "oldnum", oldNum, "new", newHead.H
        return
    }
    // If the reorg ended up on a lower number, it's indicative of seth
    log.Debug("Skipping transaction reset caused by setHead",
        "old", oldHead.Hash(), "oldnum", oldNum, "new", newHead.H
    // We still need to update the current state s.th. the lost transact
} else {
    if add == nil {
        // if the new head is nil, it means that something happened be
        // the firing of newhead-event and _now_: most likely a
        // reorg caused by sync-reversion or explicit sethead back to
        // earlier block.
        log.Warn("Transaction pool reset with missing new head", "n
        return
    }
    var discarded, included types.Transactions
    for rem.NumberU64() > add.NumberU64() {
        discarded = append(discarded, rem.Transactions()...)
        if rem = pool.chain.GetBlock(rem.ParentHash(), rem.NumberU64())
        log.Error("Unrooted old chain seen by tx pool", "block", old
        return
    }
}
for add.NumberU64() > rem.NumberU64() {
    included = append(included, add.Transactions()...)
}

```

```

        if add = pool.chain.GetBlock(add.ParentHash(), add.NumberL
        log.Error("Unrooted new chain seen by tx pool", "block", ne
        return
    }
}
for rem.Hash() != add.Hash() {
    discarded = append(discarded, rem.Transactions()...)
    if rem = pool.chain.GetBlock(rem.ParentHash(), rem.NumberL
    log.Error("Unrooted old chain seen by tx pool", "block", old
    return
}
included = append(included, add.Transactions()...)
if add = pool.chain.GetBlock(add.ParentHash(), add.NumberL
log.Error("Unrooted new chain seen by tx pool", "block", ne
return
}
}
lost := make([]*types.Transaction, 0, len(discarded))
for _, tx := range types.TxDifference(discarded, included) {
    if pool.Filter(tx) {
        lost = append(lost, tx)
    }
}
reinject = lost
}
}
// Initialize the internal state to the current head
if newHead == nil {
    newHead = pool.chain.CurrentBlock() // Special case during testing
}
statedb, err := pool.chain.StateAt(newHead.Root)
if err != nil {
    log.Error("Failed to reset txpool state", "err", err)
    return
}
}

```

```

pool.currentHead.Store(newHead)
pool.currentState = statedb
pool.pendingNonces = newNoncer(statedb)

costFn := types.NewL1CostFunc(pool.chainconfig, statedb)
pool.l1CostFn = func(dataGas types.RollupGasData) *big.Int {
    return costFn(newHead.Number.Uint64(), dataGas, false)
}

// Inject any transactions discarded due to reorgs
log.Debug("Reinjecting stale transactions", "count", len(reinject))
core.SenderCacher.Recover(pool.signer, reinject)
pool.addTxsLocked(reinject, false)
}

```

- It includes all transactions in the orphan blocks in `reinject` array, and calls `addTxsLocked` to add them inside the mempool.

`core/txpool/txpool.go`

- So how is a mempool created? `txpool.go` provides a method called `New` that allows a client to start the mempool.

#### ▼ code

- `New`

```

func New(gasTip *big.Int, chain BlockChain, subpools []SubPool) (*Tx
    // Retrieve the current head so that all subpools and this main coord
    // pool will have the same starting state, even if the chain moves for
    // during initialization.
    head := chain.CurrentBlock()

    pool := &TxPool{
        subpools:  subpools,
        reservations: make(map[common.Address]SubPool),

```

```

quit:    make(chan chan error),
}
for i, subpool := range subpools {
    if err := subpool.Init(gasTip, head, pool.reserver(i, subpool)); err != nil {
        for j := i - 1; j >= 0; j-- {
            subpools[j].Close()
        }
        return nil, err
    }
}
go pool.loop(head, chain)
return pool, nil
}

```

#### **loop** function

```

func (p *TxPool) loop(head *types.Header, chain Blockchain) {
    // Subscribe to chain head events to trigger subpool resets
    var (
        newHeadCh = make(chan core.ChainHeadEvent)
        newHeadSub = chain.SubscribeChainHeadEvent(newHeadCh)
    )
    defer newHeadSub.Unsubscribe()

    // Track the previous and current head to feed to an idle reset
    var (
        oldHead = head
        newHead = oldHead
    )
    // Consume chain head events and start resets when none is running
    var (
        resetBusy = make(chan struct{}, 1) // Allow 1 reset to run concurrently
        resetDone = make(chan *types.Header)
    )
    var errc chan error
    for errc == nil {

```

```

// Something interesting might have happened, run a reset if there
// one needed but none is running. The resetter will run on its own
// goroutine to allow chain head events to be consumed contiguous
if newHead != oldHead {
    // Try to inject a busy marker and start a reset if successful
    select {
    case resetBusy ← struct{}{}:
        // Busy marker injected, start a new subpool reset
        go func(oldHead, newHead *types.Header) {
            for _, subpool := range p.subpools {
                subpool.Reset(oldHead, newHead)
            }
            resetDone ← newHead
        }(oldHead, newHead)

    default:
        // Reset already running, wait until it finishes
    }
}
// Wait for the next chain head event or a previous reset finish
select {
case event := ←newHeadCh:
    // Chain moved forward, store the head for later consumption
    newHead = event.Block.Header()

case head := ←resetDone:
    // Previous reset finished, update the old head and allow a new
    oldHead = head
    ←resetBusy

case errc = ←p.quit:
    // Termination requested, break out on the next loop round
}
}
// Notify the closer of termination (no error possible for now)

```

```
errc ← nil  
}
```

- There is a goroutine inside the `loop` function, which subscribes to the block header change with `Reset` function.
- So if chain reorg happens, the client will include all txs in the orphan blocks.